
Essex Summer School 2023, Quantitative Text Analysis

- Instructor: dr Martijn Schoonvelde
 - martijn.schoonvelde@rug.nl
 - Office hours: by appointment (in person or via Zoom)
 - Meetings: Daily 10:00am–1:30pm (BST)
-

Course introduction

With the massive availability of text data on the web, social scientists increasingly recognize quantitative text analysis (or “text as data”) as a promising approach for analyzing various kinds of social and political phenomena. This course introduces participants to a variety of its methods and tools. We discuss their underlying theoretical assumptions, substantive applications of these methods, and their implementation in the R statistical programming language. The meetings – which combine lectures and coding sessions – will be hands-on, dealing with practical issues in each step of the research process.

Learning Outcomes

Participants will understand fundamental issues in quantitative text analysis research design such as inter-coder agreement, reliability, validation, accuracy, and precision. Participants will learn to convert texts into informative feature matrices and to analyze those matrices using statistical methods. Participants will learn to apply these methods to a text corpus in support of a substantive research question. Furthermore, participants will be able to critically evaluate (social science) research that relies on quantitative text analysis methods.

Participation and communication

I expect that you come to our meetings prepared, having read required papers, and ready to discuss your questions, criticisms and thoughts. To facilitate communication and interaction we will make use of a dedicated Slack channel at <https://essqta23.slack.com> for which I will send you an invitation via email.

Literature

We do not use a textbook in this course, but we will papers from political science and other social sciences. For further background on quantitative text analysis and natural language processing I recommend the following books:

- Daniel Jurafsky and James H. Martin (2020). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*. 3rd edition. <https://web.stanford.edu/~jurafsky/slp3/>

- Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze (2008). An Introduction to Information Retrieval. New York: Cambridge University Press. <https://nlp.stanford.edu/IR-book/information-retrieval-book.html>
- Grimmer, J., Roberts, M.E. and Stewart, B.M. (2022). Text as data: A new framework for machine learning and the social sciences. Princeton University Press.
- van Atteveldt, W., Trilling, D. and Calderon, C.A., (2022). Computational Analysis of Communication. John Wiley & Sons. <https://cssbook.net/>

Software

In this module we will use R. Students will need to have R and RStudio installed on their computers / laptops. Students who have not used R at all are advised to work their way through one of the free resources that are listed on <https://www.rstudio.com/online-learning/#R>. Another good resource is Quantitative Politics with R, developed by Erik Gahner Larsen and Zoltán Fazékas: <http://qpplr.com/index.html>. Another fantastic resource is R for Data Science by Hadley Wickham and Garrett Grolemund. This book is available at <https://r4ds.had.co.nz/>.

Course outline

** This outline serves a general plan for the course; deviations (announced) may be necessary. To keep the workload manageable we'll stick to two readings a day but during our meetings we'll discuss other papers as well.*

Day 1 – 10 July

- What is quantitative text analysis? What will you learn in this course? Developing a corpus.
 - **Required reading:**
 - * Benoit (2020). Text as Data: An Overview. Handbook of Research Methods in Political Science and International Relations. Ed. by L. Curini and R. Franzese. Thousand Oaks: Sage: 461–497.
 - * Grimmer, J., & Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, 21(3), 267–297.

Day 2 – 11 July

- Core assumptions in quantitative text analysis. Representations of text. Preprocessing and feature selection.
 - **Required reading:**
 - * Welbers, K., Van Atteveldt, W., & Benoit, K. (2017). Text analysis in R. *Communication Methods and Measures*, 11(4), 245–265.
 - * Baden, C., Pipal, C., Schoonvelde, M. & van der Velden, M.A.G., (2022). Three Gaps in Computational Text Analysis Methods for Social Sciences: A Research Agenda. *Communication Methods and Measures*, 16(1): pp. 1–18.

Day 3 – 12 July

- Advanced text representations. Risks of feature selection with unsupervised models.
 - **Required reading:**
 - * Denny, M.J. & Spirling, A., (2018). Text preprocessing for unsupervised learning: Why it matters, when it misleads, and what to do about it. *Political Analysis*, 26(2): pp.168–189.
 - * Benoit, K., Watanabe, K., Wang, H, Nulty, P., Obeng, A., Müller, & Matsuo, A. (2018). Quanteda: An R package for the quantitative analysis of textual data. *Journal of Open Source Software*, 3(30), 774.

Day 4 – 13 July

- Comparing documents in a corpus. Combining linguistic features and social science theories.
 - **Required reading:**
 - * Peterson, A. & Spirling, A., (2018). Classification accuracy as a substantive quantity of interest: Measuring polarization in Westminster systems. *Political Analysis*, 26(1): pp. 120–128.
 - * Cross, J. & Hermansson, H., (2017). Legislative amendments and informal politics in the European Union: A text reuse approach. *European Union Politics*, 18(4): pp. 581–602.

Day 5 – 14 July

- What are dictionaries and how can we validate them? Sensitivity and specificity.
 - **Required reading:**
 - * Rauh, C., 2018. Validating a sentiment dictionary for German political language—a workbench note. *Journal of Information Technology & Politics*, 15(4): pp.319–343.
 - * S.O. Proksch, W. Lowe, J. Wäckerle, & S. N. Soroka (2019). Multilingual Sentiment Analysis: A New Approach to Measuring Conflict in Legislative Speeches. *Legislative Studies Quarterly* 44(1): pp. 97–131.

Day 6 – 17 July

- Human coding and document classification using supervised machine learning. Evaluating a classifier.
 - **Required reading:**
 - * Daniel Jurafsky and James H. Martin (2020). Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition. 3rd edition: Chapter 4
 - * Müller, S., (2020). “Media coverage of campaign promises throughout the electoral cycle.” *Political Communication*: 1–23.

Day 7 – 18 July

- Supervised, semi-supervised and unsupervised approaches to place text on an underlying dimension.
 - **Required reading:**

- * Slapin J. & Proksch S. (2008). A scaling model for estimating time-serial positions from texts. *American Journal of Political Science* 52, 705–722.
- * Watanabe, K., 2021. Latent semantic scaling: A semisupervised text analysis technique for new domains and languages. *Communication Methods and Measures*, 15(2), pp.81-102.

Day 8 – 19 July

- Understanding topic models. Discussing their pros and cons?
 - **Required reading:**
 - * Blei, D. M. (2012). Probabilistic topic models. *Communications of the ACM*, 55(4), 77–84.
 - * Roberts, M et al. (2014). Structural topic models for open-ended survey responses. *American Journal of Political Science*, 58(4), 1064–1082.

Day 9 – 20 July

- New developments in data. Machine translation. Automated speech recognition. Images as data.
 - **Required reading:**
 - * Proksch, S.O., Wratil, C. and Wäckerle, J., (2019). Testing the validity of automatic speech recognition for political text analysis. *Political Analysis*, 1–21
 - * De Vries, E., Schoonvelde, M. & Schumacher, G., (2018). No longer lost in translation: Evidence that Google Translate works for comparative bag-of-words text applications. *Political Analysis*, 26(4), 417–430.
 - * Casas, A. & Williams, N.W., (2019). Images that matter: Online protests and the mobilizing role of pictures. *Political Research Quarterly*, 72(2): 360–375.

Day 10 – 21 July

- Word embeddings. Transformer-based text classification. Concluding remarks.
 - Rodman, E., (2020). A Timely Intervention: Tracking the Changing Meanings of Political Concepts with Word Vectors. *Political Analysis*, 28(1): pp. 87–111.
 - Rodriguez, P.L. and Spirling, A., (2022). Word embeddings: What works, what doesn't, and how to tell the difference for applied research. *The Journal of Politics*, 84(1): pp.101–115.
 - Chan, C.H., (2023). grafzahl: fine-tuning Transformers for text data from within R. *Computational Communication Research*, 5(1) p.76–84.